

Lab 7: Robustness Metrics

CEVE 421/521

2026-03-06

Draft Material — This content is under development and subject to change.

1 Overview

In Lab 6 you computed the **value of information** for a homeowner deciding how high to elevate their house. That analysis used two SLR scenarios (RCP 2.6 and RCP 8.5) and assumed a 50/50 prior.

But what if we don't trust those probabilities? What if someone asks: "will this decision still work if we're wrong about the future?" That's a question about **robustness**.

Today you'll evaluate the house elevation decision against **10,000 sea-level rise trajectories** spanning four physical models and four emissions pathways. You'll compute robustness metrics that don't require probabilities, discover that these metrics are sensitive to *which* scenarios you include, and then apply a principled weighting approach to make the subjective assumptions transparent.

References:

- J. D. Herman, P. M. Reed, H. B. Zeff, and G. W. Characklis [1] — Robustness metrics taxonomy
- J. Doss-Gollin and K. Keller [2] — Subjective Bayesian framework for synthesizing deep uncertainties
- K. L. Ruckert, V. Srikrishnan, and K. Keller [3] — BRICK sea-level rise projections

2 Before Lab

Complete these steps BEFORE coming to lab:

1. **Accept the GitHub Classroom assignment** (link on Canvas)
2. **Clone your repository:**

```
git clone https://github.com/CEVE-421-521/lab-07-S26-yourusername.git
cd lab-07-S26-yourusername
```

3. **Open the notebook** in VS Code or Quarto preview
4. **Run the first code cell** — it will install any missing packages (may take 10-15 minutes the first time)

! Important

Submission: You will render this notebook to PDF and submit the PDF to Canvas. See Section 8 for details.

i Note

Computation note: This lab uses a pre-built EAD emulator, so evaluating all 10,000 trajectories across all elevation heights takes only a few seconds.

3 Setup

Run the cells in this section to load packages, build the house model, and load the SLR data. You don't need to modify any of this code — just run it and read the comments to understand what it does.

3.1 Packages

```
if !isfile("Manifest.toml")
    import Pkg
    Pkg.instantiate()
end

using CairoMakie
using CSV
using DataFrames
using Distributions
using Interpolations
using NetCDF
using Random
using Statistics
using YAXArrays
using DimensionalData: dims

include("house_elevation.jl")
```

3.2 House and Surge Model

Same house and storm surge distribution as Lab 6: a single-story home in Norfolk, VA with no basement, valued at \$250,000, sitting 8 ft above the tide gauge. We hold the storm surge distribution fixed (GEV with $\mu = 3$, $\sigma = 1$, $\xi = 0.15$) so we can isolate the effect of SLR uncertainty.

```
haz_fl_dept = CSV.read("data/haz_fl_dept.csv", DataFrame)
ddf_row = filter(r -> r.DmgFnId == 105, haz_fl_dept)[1, :]
```

```

house = House(
    ddf_row;
    value_usd=250_000,
    area_ft2=1_500,
    height_above_gauge_ft=8.0,
)

surge_dist = GeneralizedExtremeValue(3.0, 1.0, 0.15)

```

3.3 EAD Emulator

Rather than recomputing the trapezoidal EAD integral for every (trajectory, year, elevation) combination, we pre-compute EAD as a function of **clearance** (house floor height above sea level) and interpolate. This turns an expensive integral into a fast table lookup.

```

ead_emulator = build_ead_emulator(house, surge_dist)

```

```

let
    clearances = range(-5, 25; length=200)
    eads = ead_emulator.(clearances)
    fig = Figure(; size=(650, 350))
    ax = Axis(fig[1, 1]; xlabel="Clearance (ft)", ylabel="Expected Annual Damage (USD)",
    title="EAD Emulator")
    lines!(ax, collect(clearances), eads; linewidth=2)
    fig
end

```




Figure 1: EAD emulator: expected annual damage as a function of clearance (house floor height above sea level). At low clearance, flooding is frequent and damages are high. At high clearance, the house is well above most storms.

3.4 SLR Trajectories

We load 10,000 sea-level rise trajectories from four physical models and four emissions pathways. The four **physical models** differ in how they represent ice sheet dynamics:

- **BRICK Slow:** Conservative ice sheet response [4]
- **BRICK Fast:** Rapid ice sheet dynamics [4]
- **K14:** Kopp et al. (2014) probabilistic projections [5]
- **DP16:** DeConto & Pollard (2016) — fast Antarctic ice sheet collapse [6]

Each model is run under four **RCP scenarios** (2.6, 4.5, 6.0, 8.5), yielding 16 model-scenario combinations with 625 trajectories each.

```
ds = open_dataset("data/slr_scenarios.nc")
```

```
const M_TO_FT = 3.28084
const MODEL_NAMES = ["BRICK_slow", "BRICK_fast", "K14", "DP16"]
const RCP_NAMES = ["rcp26", "rcp45", "rcp60", "rcp85"]

slr_m = Array(ds.slr[:, :])          # (time, trajectory) in meters
model_ids = Int.(Array(ds.model_id[:]))
rcp_ids = Int.(Array(ds.rcp_id[:]))
slr_years = Int.(Array(ds.time[:])) # 2026:2125
```

```

n_traj = size(slr_m, 2)
println("Loaded $n_traj SLR trajectories, $(length(slr_years)) years")

let
    fig = Figure(; size=(750, 400))
    ax = Axis(fig[1, 1]; xlabel="Year", ylabel="Sea Level Rise (ft)", title="SLR
Ensemble: 4 Models × 4 RCPs")
    colors = [:blue, :red, :green, :purple]

    for mid in 1:4
        mask = model_ids .== mid
        trajs_ft = slr_m[:, mask] .* M_TO_FT
        # Plot a random subset for readability
        n_plot = min(200, size(trajs_ft, 2))
        idxs = randperm(size(trajs_ft, 2))[1:n_plot]
        for j in idxs
            lines!(ax, slr_years, trajs_ft[:, j]; color=(colors[mid], 0.05), linewidth=0.5)
        end
        # Dummy line for legend (use NaN to avoid affecting axis limits)
        lines!(ax, [NaN], [NaN]; color=colors[mid], linewidth=2, label=MODEL_NAMES[mid])
    end
    axislegend(ax; position=:lt)
    fig
end

```




Figure 2: 10,000 SLR trajectories at Sewells Point, VA. Color indicates physical model. The spread reflects deep uncertainty in both emissions and ice sheet dynamics.

3.5 Performance Table

The function below computes the **total cost** (investment + discounted lifetime damages) for a given elevation height across all 10,000 SLR trajectories. We then sweep over a grid of elevation heights (0 to 14 ft) and stack the results into a single `DataFrame` called `perf`.

Each row of `perf` represents one (elevation, trajectory) combination with columns:

- `:elevation_ft` — how high we elevate (0–14 ft)
- `:model` and `:rcp` — which physical model and emissions scenario
- `:trajectory` — trajectory index (1–10,000)
- `:investment` — one-time construction cost
- `:expected_damage` — discounted lifetime damages under this SLR trajectory
- `:total_cost` — investment + expected damage (the number we want to minimize)

```
const DISCOUNT_RATE = 0.04
const PLAN_START = 2026
const PLAN_END = 2100
const PLAN_IDX = findfirst==(PLAN_START), slr_years):findfirst==(PLAN_END),
slr_years)
const PLAN_YEARS = slr_years[PLAN_IDX]
```

```
"""
Compute performance for one elevation height across all trajectories.
Returns a DataFrame with one row per trajectory.
"""

function calculate_performance(h_ft, house, slr_m, model_ids, rcp_ids)
    investment = elevation_cost(house, h_ft)
    n = size(slr_m, 2)
    results = Vector{NamedTuple}(undef, n)

    for j in 1:n
        led = 0.0
        for (i, year) in enumerate(PLAN_YEARS)
            slr_ft = slr_m[PLAN_IDX[i], j] * M_TO_FT
            clearance = house.height_above_gauge_ft + h_ft - slr_ft
            ead = ead_emulator(clearance)
            led += ead / (1.0 + DISCOUNT_RATE)^(year - PLAN_START)
        end
        results[j] = (
            elevation_ft=h_ft,
            model=MODEL_NAMES[model_ids[j]],
            rcp=RCP_NAMES[rcp_ids[j]],
            trajectory=j,
            investment=investment,
            expected_damage=led,
            total_cost=investment + led,
```

```

    )
  end
  return DataFrame(results)
end

```

```

Main.Notebook.calculate_performance

```

```

height_grid = 0:1:14
perf = vcat([calculate_performance(h, house, slr_m, model_ids, rcp_ids) for h in
height_grid]...)
println("Performance table: $(nrow(perf)) rows ($(length(height_grid)) heights ×
$n_traj trajectories)")

```

4 Act 1: Unweighted Robustness

With the `perf` table in hand, we can compute robustness metrics that treat every trajectory equally. We'll compute three metrics from J. D. Herman, P. M. Reed, H. B. Zeff, and G. W. Characklis [1]. All three answer the question “which elevation is best?” — but they define “best” differently.

4.1 Expected Cost

The simplest metric: the mean total cost across all trajectories.

$$\bar{C}(h) = \frac{1}{N} \sum_{j=1}^N C(h, s_j)$$

The code below uses `groupby` to split `perf` by elevation height, then `combine` to compute the mean of `:total_cost` within each group. This is the DataFrames equivalent of a SQL `GROUP BY` with an aggregate function.

```

expected_costs = combine(
  groupby(perf, :elevation_ft),
  :total_cost => mean => :expected_cost,
)

```

Row	elevation_ft	expected_cost
	Int64	Float64
1	0	75709.9
2	1	1.84259e5
3	2	1.66994e5
4	3	1.57834e5
5	4	152966.0
6	5	1.50381e5

Row	elevation_ft	expected_cost
	Int64	Float64
7	6	150020.0
8	7	1.50525e5
9	8	1.5116e5
10	9	1.55258e5
11	10	1.62474e5
12	11	1.69843e5
13	12	1.77268e5
14	13	1.84684e5
15	14	1.92092e5

4.2 Satisficing

What fraction of trajectories keep total cost below a threshold τ ?

$$R_{\text{sat}}(h; \tau) = \frac{1}{N} \sum_{j=1}^N \mathbb{1}[C(h, s_j) < \tau]$$

Higher is better — we want a decision that “works” in as many futures as possible.

The function below has the same structure as `expected_costs` above, but instead of `mean(costs)` you need to compute the **fraction of costs below the threshold**.

```
τ = 150_000 # cost threshold (USD)
```

```
150000
```

The function below computes satisficing, but it has a **bug**: it always returns `0.0` instead of the actual fraction. **Your task**: fix the line marked **FIX ME** so that `frac` equals the fraction of costs below τ . The function runs as-is (so the notebook renders), but the satisficing plot will be flat at zero until you fix it.

```
function satisficing(df, τ)
    combine(groupby(df, :elevation_ft)) do gdf
        costs = gdf.total_cost
        frac = 0.0 # FIX ME: compute the fraction of costs below τ
        (fraction_satisfactory=frac,)
    end
end
```

```
satisficing (generic function with 1 method)
```


i Hint

`costs <= tau` gives a vector of `true`/`false`. Taking the `mean` of a Boolean vector gives the fraction that are `true`. So: `frac = mean(costs <= tau)`.

```
sat_results = satisficing(perf, tau)
```

Row	elevation_ft	fraction_satisfactory
	Int64	Float64
1	0	0.0
2	1	0.0
3	2	0.0
4	3	0.0
5	4	0.0
6	5	0.0
7	6	0.0
8	7	0.0
9	8	0.0
10	9	0.0
11	10	0.0
12	11	0.0
13	12	0.0
14	13	0.0
15	14	0.0

i Hint

`costs <= tau` gives a vector of `true`/`false`. Taking the `mean` of a Boolean vector gives the fraction that are `true`.

4.3 Minimax Regret

Regret measures how much worse your choice is compared to the best possible choice *in that particular future*. For each trajectory j , the best possible cost is $C^*(s_j) = \min_{h'} C(h', s_j)$.

$$\text{Regret}(h, s_j) = C(h, s_j) - C^*(s_j)$$

Minimax regret finds the elevation that minimizes the worst-case regret:

$$R_{\text{regret}}(h) = \max_j \text{Regret}(h, s_j)$$

The code below computes this in three steps:

1. For each trajectory, find the minimum cost across all elevation heights (the “best you could have done” in that future)
2. Join that back to `perf` and compute regret as the difference
3. For each elevation, take the maximum regret across all trajectories

```
# Step 1: best possible cost in each trajectory
best_per_traj = combine(groupby(perf, :trajectory), :total_cost => minimum
=> :best_cost)

# Step 2: join back and compute regret
perf_with_regret = leftjoin(perf, best_per_traj; on=:trajectory)
perf_with_regret.regret = perf_with_regret.total_cost .- perf_with_regret.best_cost

# Step 3: worst-case regret for each elevation
regret_results = combine(
    groupby(perf_with_regret, :elevation_ft),
    :regret => maximum => :max_regret,
)
```

Row	elevation_ft	max_regret
	Int64	Float64
1	0	2.31887e5
2	1	2.64408e5
3	2	1.84533e5
4	3	125601.0
5	4	1.06312e5
6	5	1.04949e5
7	6	1.05162e5
8	7	1.06191e5
9	8	1.06901e5
10	9	1.11328e5
11	10	1.18597e5
12	11	1.2606e5
13	12	1.33493e5
14	13	1.409e5
15	14	1.48302e5

4.4 Comparing the Three Metrics

```
let
  fig = Figure(; size=(900, 350))
  hg = collect(height_grid)

  # Expected cost
  ax1 = Axis(fig[1, 1]; xlabel="Elevation (ft)", ylabel="Expected Cost (USD)",
title="Expected Cost")
  ec = expected_costs.expected_cost
  lines!(ax1, hg, ec; linewidth=2, color=:black)
  idx = argmin(ec)
  scatter!(ax1, [hg[idx]], [ec[idx]]; color=:black, markersize=15, marker=:star5)

  # Satisficing
  ax2 = Axis(fig[1, 2]; xlabel="Elevation (ft)", ylabel="Fraction Satisfactory",
title="Satisficing ( $\tau = \frac{\tau}{1000}K$ )")
  sf = sat_results.fraction_satisfactory
  lines!(ax2, hg, sf; linewidth=2, color=:teal)
  idx = argmax(sf)
  scatter!(ax2, [hg[idx]], [sf[idx]]; color=:teal, markersize=15, marker=:star5)

  # Minimax regret
  ax3 = Axis(fig[1, 3]; xlabel="Elevation (ft)", ylabel="Max Regret (USD)",
title="Minimax Regret")
  mr = regret_results.max_regret
  lines!(ax3, hg, mr; linewidth=2, color=:firebrick)
  idx = argmin(mr)
  scatter!(ax3, [hg[idx]], [mr[idx]]; color=:firebrick, markersize=15, marker=:star5)

  fig
end
```




Figure 3: Three robustness metrics as a function of elevation height. Left: expected cost (lower is better). Center: satisficing at \$150K threshold (higher is better). Right: minimax regret (lower is better). Stars mark the best elevation under each metric.

Your Response — Question 1

Do the three metrics agree on the best elevation height? If they disagree, explain intuitively why each metric favors a different choice. Which metric would you use if advising a homeowner? Why?

5 Act 2: Sensitivity to Scenario Selection

The robustness metrics above treated all 10,000 trajectories equally. But the set of trajectories we included was itself a choice — we picked 4 models and 4 RCPs. What happens if we change that choice?

No new simulations needed — we just `filter` the `perf` DataFrame to keep only certain models or RCPs, then recompute the same metrics on each subset. The function `compute_metrics` wraps up the three metric computations from Act 1 so we can apply them to different subsets easily.

```
function compute_metrics(df, τ)
  ec = combine(groupby(df, :elevation_ft), :total_cost => mean => :expected_cost)

  sat = combine(groupby(df, :elevation_ft), :total_cost => (c -> mean(c .< τ))
=> :fraction_satisfactory)

  best = combine(groupby(df, :trajectory), :total_cost => minimum => :best_cost)
  df_r = leftjoin(df, best; on=:trajectory)
  df_r.regret = df_r.total_cost .- df_r.best_cost
  reg = combine(groupby(df_r, :elevation_ft), :regret => maximum => :max_regret)

  return (expected_cost=ec, satisficing=sat, regret=reg)
end
```

```
compute_metrics (generic function with 1 method)
```

```
subsets = [
  ("All scenarios", perf),
  ("Drop RCP 8.5", filter(r -> r.rcp != "rcp85", perf)),
  ("Drop RCP 2.6", filter(r -> r.rcp != "rcp26", perf)),
  ("Drop DP16", filter(r -> r.model != "DP16", perf)),
  ("BRICK Slow + RCP 4.5 only", filter(r -> r.model == "BRICK_slow" && r.rcp ==
"rcp45", perf)),
]
```

```
5-element Vector{Tuple{String, DataFrame}}:
```

```
("All scenarios", 150000x7 DataFrame
```

Row	elevation_ft	model	rcp	trajectory	investment	expected_d ...
	Int64	String	String	Int64	Float64	Float64 ...
1	0	BRICK_fast	rcp26	1	0.0	5981 ...
2	0	BRICK_fast	rcp26	2	0.0	7405
3	0	BRICK_fast	rcp26	3	0.0	6249
4	0	BRICK_fast	rcp26	4	0.0	6054
5	0	BRICK_fast	rcp26	5	0.0	6012 ...
6	0	BRICK_fast	rcp26	6	0.0	5695
7	0	BRICK_fast	rcp26	7	0.0	6170
8	0	BRICK_fast	rcp26	8	0.0	6132
⋮	⋮	⋮	⋮	⋮	⋮	⋮
149994	14	K14	rcp85	9994	191370.0	74 ...
149995	14	K14	rcp85	9995	191370.0	72
149996	14	K14	rcp85	9996	191370.0	72
149997	14	K14	rcp85	9997	191370.0	73
149998	14	K14	rcp85	9998	191370.0	71 ...
149999	14	K14	rcp85	9999	191370.0	71
150000	14	K14	rcp85	10000	191370.0	69

```
2 columns and 149985 rows omitted)
```

```
("Drop RCP 8.5", 112500x7 DataFrame
```

Row	elevation_ft	model	rcp	trajectory	investment	expected_d ...
	Int64	String	String	Int64	Float64	Float64 ...
1	0	BRICK_fast	rcp26	1	0.0	5981 ...
2	0	BRICK_fast	rcp26	2	0.0	7405
3	0	BRICK_fast	rcp26	3	0.0	6249
4	0	BRICK_fast	rcp26	4	0.0	6054
5	0	BRICK_fast	rcp26	5	0.0	6012 ...
6	0	BRICK_fast	rcp26	6	0.0	5695
7	0	BRICK_fast	rcp26	7	0.0	6170
8	0	BRICK_fast	rcp26	8	0.0	6132

:	:	:	:	:	:	:	:	:
112494		14	K14	rcp60	9369	191370.0	71	...
112495		14	K14	rcp60	9370	191370.0	67	
112496		14	K14	rcp60	9371	191370.0	69	
112497		14	K14	rcp60	9372	191370.0	74	
112498		14	K14	rcp60	9373	191370.0	70	...
112499		14	K14	rcp60	9374	191370.0	74	
112500		14	K14	rcp60	9375	191370.0	71	

2 columns and 112485 rows omitted)

("Drop RCP 2.6", 112500x7 DataFrame

Row	elevation_ft	model	rcp	trajectory	investment	expected_d	...
	Int64	String	String	Int64	Float64	Float64	...

1	0	BRICK_fast	rcp45	626	0.0	7126	...
2	0	BRICK_fast	rcp45	627	0.0	6644	
3	0	BRICK_fast	rcp45	628	0.0	6648	
4	0	BRICK_fast	rcp45	629	0.0	9261	
5	0	BRICK_fast	rcp45	630	0.0	6136	...
6	0	BRICK_fast	rcp45	631	0.0	6774	
7	0	BRICK_fast	rcp45	632	0.0	6224	
8	0	BRICK_fast	rcp45	633	0.0	6259	

:	:	:	:	:	:	:	:	
112494		14	K14	rcp85	9994	191370.0	74	...
112495		14	K14	rcp85	9995	191370.0	72	
112496		14	K14	rcp85	9996	191370.0	72	
112497		14	K14	rcp85	9997	191370.0	73	
112498		14	K14	rcp85	9998	191370.0	71	...
112499		14	K14	rcp85	9999	191370.0	71	
112500		14	K14	rcp85	10000	191370.0	69	

2 columns and 112485 rows omitted)

("Drop DP16", 112500x7 DataFrame

Row	elevation_ft	model	rcp	trajectory	investment	expected_d	...
	Int64	String	String	Int64	Float64	Float64	...

1	0	BRICK_fast	rcp26	1	0.0	5981	...
2	0	BRICK_fast	rcp26	2	0.0	7405	
3	0	BRICK_fast	rcp26	3	0.0	6249	
4	0	BRICK_fast	rcp26	4	0.0	6054	
5	0	BRICK_fast	rcp26	5	0.0	6012	...
6	0	BRICK_fast	rcp26	6	0.0	5695	
7	0	BRICK_fast	rcp26	7	0.0	6170	
8	0	BRICK_fast	rcp26	8	0.0	6132	

:	:	:	:	:	:	:	:	
112494		14	K14	rcp85	9994	191370.0	74	...
112495		14	K14	rcp85	9995	191370.0	72	
112496		14	K14	rcp85	9996	191370.0	72	

112497	14	K14	rcp85	9997	191370.0	73
112498	14	K14	rcp85	9998	191370.0	71 ...
112499	14	K14	rcp85	9999	191370.0	71
112500	14	K14	rcp85	10000	191370.0	69

2 columns and 112485 rows omitted)

("BRICK Slow + RCP 4.5 only", 9375x7 DataFrame

Row	elevation_ft	model	rcp	trajectory	investment	expected_dam	...
	Int64	String	String	Int64	Float64	Float64	...
1	0	BRICK_slow	rcp45	3126	0.0	61726.	...
2	0	BRICK_slow	rcp45	3127	0.0	60408.	
3	0	BRICK_slow	rcp45	3128	0.0	65041.	
4	0	BRICK_slow	rcp45	3129	0.0	66624.	
5	0	BRICK_slow	rcp45	3130	0.0	70243.	...
6	0	BRICK_slow	rcp45	3131	0.0	60148.	
7	0	BRICK_slow	rcp45	3132	0.0	61740.	
8	0	BRICK_slow	rcp45	3133	0.0	60651.	
:	:	:	:	:	:	:	...
9369	14	BRICK_slow	rcp45	3744	191370.0	707.	...
9370	14	BRICK_slow	rcp45	3745	191370.0	702.	
9371	14	BRICK_slow	rcp45	3746	191370.0	705.	
9372	14	BRICK_slow	rcp45	3747	191370.0	716.	
9373	14	BRICK_slow	rcp45	3748	191370.0	712.	...
9374	14	BRICK_slow	rcp45	3749	191370.0	715.	
9375	14	BRICK_slow	rcp45	3750	191370.0	710.	

2 columns and 9360 rows omitted)

```
let
    fig = Figure(; size=(900, 350))
    hg = collect(height_grid)
    colors = Makie.wong_colors()

    ax1 = Axis(fig[1, 1]; xlabel="Elevation (ft)", ylabel="Expected Cost (USD)",
    title="Expected Cost")
    ax2 = Axis(fig[1, 2]; xlabel="Elevation (ft)", ylabel="Fraction Satisfactory",
    title="Satisficing ( $\tau = \frac{\tau}{1000}K$ )")
    ax3 = Axis(fig[1, 3]; xlabel="Elevation (ft)", ylabel="Max Regret (USD)",
    title="Minimax Regret")

    for (i, (label, df)) in enumerate(subsets)
        m = compute_metrics(df,  $\tau$ )
        lines!(ax1, hg, m.expected_cost.expected_cost; linewidth=2, color=colors[i],
        label=label)
        lines!(ax2, hg, m.satisficing.fraction_satisfactory; linewidth=2,
        color=colors[i], label=label)
```

```

        lines!(ax3, hg, m.regret.max_regret; linewidth=2, color=colors[i], label=label)
    end
    Legend(fig[2, :], ax1; orientation=:horizontal, tellwidth=false, tellheight=true)
    fig
end

```




Figure 4: Sensitivity of robustness metrics to scenario selection. Each color represents a different subset of the SLR ensemble. The ‘best’ elevation shifts depending on which futures are included — even though no explicit probabilities were assigned.

Your Response — Question 2

How much do the robustness metrics change when you drop RCP 8.5? When you drop DP16? What does this tell you about the relationship between “robustness” and the choice of scenario ensemble? Is it defensible to claim a decision is “robust” without specifying which futures were considered?

6 Act 3: Weighted Robustness

J. Doss-Gollin and K. Keller [2] propose a framework for making subjective beliefs about SLR explicit. Instead of treating all trajectories equally, we assign **weights** based on a prior belief about how much sea level will rise.

6.1 Weighting Method and Priors

We project each trajectory onto a single number: SLR at the end of the planning horizon (2100). Then we define a probability distribution $p_{\text{belief}}(\psi)$ over that value and compute weights:

$$w_j = F_{\text{belief}}\left(\frac{\psi_j + \psi_{j+1}}{2}\right) - F_{\text{belief}}\left(\frac{\psi_{j-1} + \psi_j}{2}\right)$$

where F_{belief} is the CDF of the prior and ψ_j are the sorted SLR values. This assigns more weight to trajectories near the center of the belief distribution and less to trajectories in the tails.

Following J. Doss-Gollin and K. Keller [2], we define three Gamma priors over SLR at Sewells Point in 2100 (in feet). These represent different beliefs about how much sea level will rise — from optimistic to pessimistic.

```
priors = [
    (name="Slow SLR", dist=Gamma(3.0, 0.5)),      # E[X] = 1.5 ft
    (name="Moderate SLR", dist=Gamma(3.0, 1.0)),  # E[X] = 3.0 ft
    (name="Rapid SLR", dist=Gamma(5.0, 1.0)),     # E[X] = 5.0 ft
]
```

The `make_weights` function below extracts SLR in 2100 for each trajectory, sorts them, computes CDF-based weights, normalizes to sum to 1, and maps the weights back to the original trajectory order. You don't need to modify this — just understand that `make_weights(prior, slr_values)` returns a vector of weights (one per trajectory) that sum to 1.

```
idx_2100 = findfirst(==(2100), slr_years)
slr_2100_ft = slr_m[idx_2100, :] .* M_TO_FT

"""Compute trajectory weights from a prior distribution over SLR in 2100 (ft)."""
function make_weights(prior_dist, slr_values)
    n = length(slr_values)
    order = sortperm(slr_values)
    ψ = slr_values[order]

    # Compute CDF at midpoints between sorted values
    w = zeros(n)
    for j in 1:n
        if j == 1
            upper = (ψ[1] + ψ[2]) / 2
            w[j] = cdf(prior_dist, upper)
        elseif j == n
            lower = (ψ[n - 1] + ψ[n]) / 2
            w[j] = 1.0 - cdf(prior_dist, lower)
        else
            upper = (ψ[j] + ψ[j + 1]) / 2
            lower = (ψ[j - 1] + ψ[j]) / 2
            w[j] = cdf(prior_dist, upper) - cdf(prior_dist, lower)
        end
    end

    # Normalize and unsort
    w ./= sum(w)
    result = similar(w)
    result[order] = w
    return result
end
```

```
Main.Notebook.make_weights
```

```
let
    fig = Figure(; size=(700, 350))
    ax = Axis(fig[1, 1]; xlabel="SLR in 2100 (ft)", ylabel="Density", title="SLR Priors
vs. Ensemble Distribution")
    hist!(ax, slr_2100_ft; bins=80, normalization=:pdf, color=(:gray, 0.3),
strokewidth=0.5, strokecolor=:gray)
    x = range(-1, 15; length=500)
    colors = [:blue, :orange, :red]
    for (i, p) in enumerate(priors)
        lines!(ax, collect(x), pdf.(Ref(p.dist), x); linewidth=2, color=colors[i],
label=p.name)
    end
    axislegend(ax; position=:rt)
    fig
end
```




Figure 5: Three prior beliefs about SLR at Sewells Point in 2100 (ft), overlaid with the histogram of the 10,000 ensemble trajectories. The priors span a range from optimistic (Slow SLR) to pessimistic (Rapid SLR).

6.2 Weighted Metrics

With weights, expected cost becomes a **weighted sum** instead of a simple mean:

$$\bar{C}_w(h) = \sum_{j=1}^N w_j \cdot C(h, s_j)$$

The two functions below have the same structure as the unweighted metrics, but they have **bugs**: they ignore the weights and just compute the unweighted versions. **Your task**: fix the lines marked **FIX ME** so that:

- `weighted_expected_cost` computes a **weighted sum** of costs instead of a simple mean
- `weighted_satisficing` computes a **weighted fraction** below the threshold instead of an unweighted fraction

The variable `w` contains the weights for this group's trajectories (already extracted for you) and `costs` contains the total costs. Both functions run as-is, but the weighted plot will look identical to the unweighted plot until you fix them.

```
function weighted_expected_cost(df, weights)
    combine(groupby(df, :elevation_ft)) do gdf
        w = weights[gdf.trajectory]
        costs = gdf.total_cost
        (expected_cost=mean(costs),) # FIX ME: use weights instead of simple mean
    end
end

function weighted_satisficing(df, weights, τ)
    combine(groupby(df, :elevation_ft)) do gdf
        w = weights[gdf.trajectory]
        costs = gdf.total_cost
        (fraction_satisfactory=mean(costs .< τ),) # FIX ME: use weights instead of
simple mean
    end
end
```

```
weighted_satisficing (generic function with 1 method)
```

i Hint

The unweighted `mean(costs)` treats every trajectory as equally likely ($1/N$ each). With weights, replace `mean(costs)` with `sum(costs .* w)`, and `mean(costs .< τ)` with `sum(w .* (costs .< τ))`.

```
let
    fig = Figure(; size=(700, 350))
    hg = collect(height_grid)
    colors = [:blue, :orange, :red]

    ax1 = Axis(fig[1, 1]; xlabel="Elevation (ft)", ylabel="Expected Cost (USD)",
title="Weighted Expected Cost")
    ax2 = Axis(fig[1, 2]; xlabel="Elevation (ft)", ylabel="Fraction Satisfactory",
title="Weighted Satisficing (τ = \$(τ÷1000)K)")
```

```

for (i, p) in enumerate(priors)
  w = make_weights(p.dist, slr_2100_ft)
  wec = weighted_expected_cost(perf, w)
  wsat = weighted_satisficing(perf, w,  $\tau$ )
  lines!(ax1, hg, wec.expected_cost; linewidth=2, color=colors[i], label=p.name)
  lines!(ax2, hg, wsat.fraction_satisfactory; linewidth=2, color=colors[i],
label=p.name)
end
Legend(fig[2, :], ax1; orientation=:horizontal, tellwidth=false, tellheight=true)
fig
end

```




Figure 6: Weighted robustness metrics under three subjective priors. Left: expected cost (lower is better). Right: satisfying (higher is better). The belief about SLR shifts the optimal elevation, but the analysis is now transparent about which assumptions drive the recommendation.

Your Response — Question 3

Compare the weighted results to the unweighted results from Act 1. How does the choice of prior change the recommended elevation? Does the weighting stabilize the results compared to the subset sensitivity in Act 2?

7 Wrap-Up

Your Response — Question 4

You are advising a city council in Norfolk, VA on setting minimum elevation requirements for new construction.

1. *Which elevation height would you recommend, and which robustness metric would you use to justify your recommendation?*
2. *A council member asks: "Is this robust?" How do you respond?*
3. *What are the trade-offs between the three approaches (expected cost, satisficing, minimax regret)?*

Key takeaways:

1. **Robustness metrics** let you evaluate decisions across many uncertain futures without requiring probabilities.
2. **No metric is "objective"**: the choice of scenarios, thresholds, and metrics all involve subjective judgments.
3. **Making assumptions explicit** (via weighted robustness) is more transparent than hiding them in the choice of scenario ensemble.
4. **Different metrics favor different decisions**: expected cost, satisficing, and minimax regret encode different attitudes toward risk and uncertainty.

8 Submission

1. Write your answers in the response boxes above.
2. **Render to PDF**:
 - In VS Code: Open the command palette (Cmd+Shift+P / Ctrl+Shift+P) → "Quarto: Render Document" → select Typst PDF
 - Or from the terminal: `quarto render index.qmd --to typst`
3. **Submit the PDF** to the Lab 7 assignment on Canvas.
4. **Push your code** to GitHub (for backup):

```
git add -A && git commit -m "Lab 7 complete" && git push
```

Checklist

Before submitting:

- ☐ Packages load without error
- ☐ EAD emulator plotted (Figure 1)
- ☐ SLR ensemble plotted (Figure 2)
- ☐ Performance table computed (all heights × all trajectories)
- ☐ Satisficing function completed (Figure 3)
- ☐ Question 1 answered (comparing metrics)

-  Scenario sensitivity plotted (Figure 4)
-  Question 2 answered (sensitivity interpretation)
-  Weighted expected cost and satisficing completed (Figure 6)
-  Question 3 answered (weighted vs. unweighted)
-  Question 4 answered (recommendation)
-  Notebook renders to PDF without errors

Bibliography

- [1] J. D. Herman, P. M. Reed, H. B. Zeff, and G. W. Characklis, “How Should Robustness Be Defined for Water Systems Planning under Change?,” *Journal of Water Resources Planning and Management*, vol. 141, no. 10, p. 4015012, Oct. 2015, doi: 10.1061/(asce)wr.1943-5452.0000509.
- [2] J. Doss-Gollin and K. Keller, “A Subjective Bayesian Framework for Synthesizing Deep Uncertainties in Climate Risk Management,” *Earth’s Future*, vol. 11, no. 1, p. e2022EF003044, Jan. 2023, doi: 10.1029/2022EF003044.
- [3] K. L. Ruckert, V. Srikrishnan, and K. Keller, “Characterizing the Deep Uncertainties Surrounding Coastal Flood Hazard Projections: A Case Study for Norfolk, VA,” *Scientific Reports*, vol. 9, no. 1, p. 11373, Aug. 2019, doi: 10.1038/s41598-019-47587-6.
- [4] T. E. Wong, A. M. R. Bakker, K. Ruckert, P. Applegate, A. B. A. Slangen, and K. Keller, “BRICK v0.2, a Simple, Accessible, and Transparent Model Framework for Climate and Regional Sea-Level Projections,” *Geoscientific Model Development*, vol. 10, no. 7, pp. 2741–2760, July 2017, doi: 10.5194/gmd-10-2741-2017.
- [5] R. E. Kopp *et al.*, “Probabilistic 21st and 22nd Century Sea-Level Projections at a Global Network of Tide-Gauge Sites,” *Earth’s Future*, vol. 2, no. 8, pp. 383–406, 2014, doi: 10.1002/2014ef000239.
- [6] R. M. DeConto and D. Pollard, “Contribution of Antarctica to Past and Future Sea-Level Rise,” *Nature*, vol. 531, no. 7596, pp. 591–597, Mar. 2016, doi: 10.1038/nature17145.