

Lab 8: Scenario Discovery

CEVE 421/521

2026-03-27

1 Overview

In this lab you'll use **scenario discovery for policy design** — asking not just “when does a fixed policy fail?” but “what combinations of uncertainty and policy choice drive high costs?” You'll work with the Eijgenraam dike heightening model, which extends D. van Dantzig [1] to consider time dependence in a cost-benefit framework for setting flood protection standards in the Netherlands [2]. You'll jointly sample uncertain climate parameters and policy levers, label each (uncertainty, policy) pair by whether total cost exceeds a threshold, and use **PRIM** (Patient Rule Induction Method) to find which combinations drive expensive outcomes [3].

2 Before Lab

Complete these steps **BEFORE** coming to lab:

1. **Accept the GitHub Classroom assignment** (link on Canvas)
2. **Clone your repository:**

```
git clone https://github.com/CEVE-421-521/lab-08-S26-yourusername.git
cd lab-08-S26-yourusername
```

3. **Open the notebook** in VS Code or Quarto preview
4. **Run the first code cell** — it will install any missing packages (may take 10-15 minutes the first time)

! Important

Submission: You will render this notebook to PDF and submit the PDF to Canvas. See Section 6 for details.

2.1 Setup

Run the cells in this section to load packages, the model, and generate the ensemble. You don't need to modify any of this code — just run it and read the comments.

We use `SimOptDecisions` as before

```

# Install SimOptDecisions if needed
if !isfile("Manifest.toml")
    import Pkg
    Pkg.instantiate()
end

try
    using SimOptDecisions
catch
    import Pkg
    try
        Pkg.rm("SimOptDecisions")
    catch
    end
    Pkg.add(; url="https://github.com/dossgollin-lab/SimOptDecisions")
    Pkg.instantiate()
    using SimOptDecisions
end

```

We also use regular packages

```

using CairoMakie
using DataFrames
using Distributions
using Random
using Statistics

```

Last, we implement PRIM and the case study in separate files

```

include("eijgenraam.jl")
include("prim.jl")

Random.seed!(2026)

```

2.2 The Eijgenraam Dike Model

The Eijgenraam model [2] computes the total societal cost of flood protection for a dike ring in the Netherlands. It balances two competing forces:

- **Investment cost:** building the dike higher is expensive, and costs grow exponentially with height
- **Expected flood losses:** without protection, losses grow over time due to sea level rise and economic growth

The annual **flood probability** at year t is:

$$P(t) = P_0 \exp(\alpha \eta t - \alpha H(t))$$

where P_0 is the initial flood probability, α controls how sensitive flood probability is to water level, η is the rate of water level rise, and $H(t)$ is the cumulative dike heightening at time t .

i Note

Notice that α appears **twice**: it makes flood probability grow faster with rising water ($\alpha\eta t$), but it also makes each cm of dike height more effective ($-\alpha H$). This means α 's effect on failure is **non-monotonic** — something scenario discovery can reveal.

2.3 Policy: When and How Much to Heighten

Our policy has two levers:

- **When** to heighten the dike (t_h years from now)
- **How much** to heighten it (u cm)

There's a genuine tradeoff: deferring the investment saves money (discounting), but the dike ring is unprotected while we wait.

2.4 Sampling Design

For this analysis we fix three parameters at calibrated Ring 15 values and jointly sample two climate uncertainties and two policy levers.

Fixed (calibrated):

- $P_0 = 1/2000$ — the current legal safety standard, used as the nominal initial flood probability
- $V_0 = 15,000$ M€ — approximate economic value at risk for Ring 15
- $\delta = 0.04$ — discount rate

Jointly sampled:

Variable	Role	Range
α	sensitivity of flood prob. to water level	[0.025, 0.075]
η	sea level rise rate (cm/year)	[0.4, 1.2]
t_{heighten}	year of dike heightening	[0, 50]
u_{heighten}	amount of heightening (cm)	[50, 250]

Each of the 1,000 runs pairs one randomly drawn uncertainty realization with one randomly drawn policy. PRIM will then search this joint space for conditions — combining climate and policy — that drive high total cost.

```
config = EijgenraamConfig() # ring 15 structural defaults -- see eijgenraam.jl

# Fixed calibrated parameters
P0_fixed = 1 / 2000
V0_fixed = 15_000.0
```

```

delta_fixed = 0.04

n_scenarios = 1000
scenarios = map(1:n_scenarios) do _
    EigenraamScenario(;
        P0=P0_fixed,
        alpha=rand(Uniform(0.025, 0.075)),
        eta=rand(Uniform(0.4, 1.2)),
        V0=V0_fixed,
        delta=delta_fixed,
    )
end

policies = map(1:n_scenarios) do _
    EigenraamPolicy(;
        t_heighten=rand(Uniform(0.0, 50.0)),
        u_heighten=rand(Uniform(50.0, 250.0)),
    )
end

```

3 Exploratory Modeling

3.1 Run the Ensemble

We run each (uncertainty, policy) pair through the model individually and collect the outcomes.

```

results_vec = map(1:n_scenarios) do i
    r = explore(config, scenarios[i:i], [policies[i]])
    (
        total_cost=Array(r.total_cost)[1, 1],
        total_investment=Array(r.total_investment)[1, 1],
        total_loss=Array(r.total_loss)[1, 1],
        max_failure_prob=Array(r.max_failure_prob)[1, 1],
    )
end

total_costs = [r.total_cost for r in results_vec]
total_investments = [r.total_investment for r in results_vec]
total_losses = [r.total_loss for r in results_vec]

```

3.2 Policy-Cost Landscape

The scatter plot below shows total cost for each run in policy space. Each point is one (uncertainty, policy) pair — the scatter around any policy reflects variation due to the sampled climate parameters.

```

let
    t_h_vals = [SimOptDecisions.value(p.t_heighten) for p in policies]
    u_vals = [SimOptDecisions.value(p.u_heighten) for p in policies]

    fig = Figure(; size=(700, 450))
    ax = Axis(
        fig[1, 1];
        xlabel="Year of heightening",
        ylabel="Heightening amount (cm)",
        title="Total Cost by Policy Choice",
    )
    sc = scatter!(ax, t_h_vals, u_vals; color=total_costs, colormap=:viridis,
markersize=5, alpha=0.7)
    Colorbar(fig[1, 2], sc; label="Total Cost (M€)")
    fig
end

```




Figure 1: Total cost for each (uncertainty, policy) pair, plotted in policy space. Points higher and to the right (large heightening, done early) tend to have higher investment cost but lower expected losses.

4 Scenario Discovery

Now we ask: **under what joint conditions of climate and policy does total cost become unacceptably high?**

4.1 Cost Budget

Rather than an arbitrary statistical threshold, we set a cost budget from an approximate analysis. We run the benchmark policy — heighten 100 cm immediately — under a reference scenario with climate parameters at the midpoints of their ranges. Any (scenario, policy) pair that exceeds this benchmark cost is labeled a failure.

```
ref_scenario = EijgenraamScenario(  
    P0=P0_fixed,  
    alpha=0.05,    # midpoint of [0.025, 0.075]  
    eta=0.8,       # midpoint of [0.4, 1.2]  
    V0=V0_fixed,  
    delta=delta_fixed,  
)  
benchmark_policy = EijgenraamPolicy(; t_heighten=0.0, u_heighten=100.0)  
ref_result = explore(config, [ref_scenario], [benchmark_policy])  
cost_budget = Array(ref_result.total_cost)[1, 1]  
println("Cost budget: $(round(cost_budget; digits=0)) M€ (benchmark policy under  
reference scenario)")
```

```
Cost budget: 672.0 M€ (benchmark policy under reference scenario)
```

i Note

This budget reflects a specific value judgment: a reasonable policy under typical conditions. A stricter budget would label more runs as failures; a looser one fewer. The choice shapes which conditions PRIM finds — we'll return to this below.

4.2 Failure Criterion

Your Response — Question 1

The cost budget above was computed for a specific benchmark policy ($u = 100$ cm, $t = 0$) and reference scenario ($\alpha = 0.05$, $\eta = 0.8$). Try changing one of these choices — use a stricter benchmark (e.g., $u = 150$ cm), a different reference climate, or define failure some other way entirely. Re-run the failure labeling and report: how many runs are now labeled failures? In 2–3 sentences, how might a different failure definition change what PRIM finds?

```
# Your code here
```

4.3 Outcome Classification

```
failures = total_costs .> cost_budget
n_fail = sum(failures)
println("$(n_fail) of $(n_scenarios) runs exceed the budget ($(round(100 * n_fail /
n_scenarios; digits=1))%)")
```

```
684 of 1000 runs exceed the budget (68.4%)
```

4.4 Input Matrix

To run PRIM, we need the sampled inputs as a plain matrix. Each row is one (uncertainty, policy) pair; each column is one input. We include both climate uncertainties and policy levers — PRIM will search across all of them.

```
param_names = ["α", "η", "t_heighten", "u_heighten"]
X = hcat(
    [SimOptDecisions.value(s.alpha) for s in scenarios],
    [SimOptDecisions.value(s.eta) for s in scenarios],
    [SimOptDecisions.value(p.t_heighten) for p in policies],
    [SimOptDecisions.value(p.u_heighten) for p in policies],
)
```

4.5 PRIM Analysis

PRIM peels away slices of the input space to find a “box” where failures concentrate. It traces a path from high coverage (large box, catches most failures) to high density (small box, almost all scenarios in the box are failures).

```
trajectory = prim_peel(X, failures)
prim_summary(trajectory, param_names)
```

Step	Coverage	Density	Support	Restricted dimensions
----	-----	-----	-----	-----
1	100.0%	68.4%	1000	(full space)
2	97.4%	70.1%	950	u_heighten ≥ 61.412
3	95.3%	72.2%	903	u_heighten ≥ 70.2828
4	93.0%	74.1%	858	u_heighten ≥ 80.3508
5	91.7%	76.8%	816	u_heighten ≥ 90.4008
6	90.9%	80.2%	776	u_heighten ≥ 98.5278
7	89.6%	83.1%	738	u_heighten ≥ 105.9284
8	88.0%	85.8%	702	u_heighten ≥ 112.6945
9	86.0%	88.2%	667	u_heighten ≥ 118.2585
10	84.2%	90.9%	634	u_heighten ≥ 123.8134
11	82.0%	93.0%	603	u_heighten ≥ 132.2375
12	78.5%	93.7%	573	t_heighten ≤ 47.0415, u_heighten ≥ 132.2375

13		75.3%		94.5%		545		$t_{\text{heighten}} \leq 43.8171, u_{\text{heighten}} \geq 132.2375$
14		72.1%		95.2%		518		$t_{\text{heighten}} \leq 43.8171, u_{\text{heighten}} \geq 138.6029$
15		69.7%		96.8%		493		$t_{\text{heighten}} \leq 43.8171, u_{\text{heighten}} \geq 142.9466$
16		67.0%		97.7%		469		$t_{\text{heighten}} \leq 43.8171, u_{\text{heighten}} \geq 148.1629$
17		64.2%		98.4%		446		$t_{\text{heighten}} \leq 43.8171, u_{\text{heighten}} \geq 154.0102$
18		61.4%		99.1%		424		$t_{\text{heighten}} \leq 43.8171, u_{\text{heighten}} \geq 159.5594$
19		58.6%		99.5%		403		$t_{\text{heighten}} \leq 43.8171, u_{\text{heighten}} \geq 164.3972$
20		55.8%		99.7%		383		$\eta \geq 0.4516, t_{\text{heighten}} \leq 43.8171, u_{\text{heighten}} \geq 164.3972$

```

let
    coverages = [b.coverage for b in trajectory]
    densities = [b.density for b in trajectory]

    fig = Figure(; size=(600, 400))
    ax = Axis(
        fig[1, 1];
        xlabel="Density (fraction of box that are failures)",
        ylabel="Coverage (fraction of all failures captured)",
        title="PRIM Peeling Trajectory",
    )
    scatterlines!(ax, densities, coverages; markersize=8, linewidth=2)

    # Label first and last
    text!(ax, densities[1], coverages[1]; text="Start", align=:right, :bottom),
    fontsize=12)
    text!(
        ax,
        densities[end],
        coverages[end];
        text="End",
        align=:left, :top),
        fontsize=12,
    )
    fig
end

```





Figure 2: PRIM peeling trajectory: each point is a box. Moving right increases density (precision); moving down decreases coverage (recall). The analyst chooses a box that balances both.

4.6 Box Selection

Pick a box from the trajectory that you think gives a good balance of coverage and density. A good starting point is a box with density above 70% and coverage above 50%.

```
# Find a box with density > 0.7 (or the last box if none qualify)
good_idx = findfirst(b -> b.density > 0.7, trajectory)
if good_idx === nothing
    good_idx = length(trajectory)
end
box = trajectory[good_idx]
println("Selected box $(good_idx):")
println(" Coverage: $(round(100 * box.coverage; digits=1))%")
println(" Density:  $(round(100 * box.density; digits=1))%")
println(" Support:  $(box.support) scenarios")
println()
initial = trajectory[1]
for (j, name) in enumerate(param_names)
    lo_changed = box.lower[j] > initial.lower[j]
    hi_changed = box.upper[j] < initial.upper[j]
    if lo_changed || hi_changed
        println(" $name: [$(round(box.lower[j]; sigdigits=3)), $(round(box.upper[j];
```

```
sigdigits=3))])")
    end
end
```

```
Selected box 2:
  Coverage: 97.4%
  Density:  70.1%
  Support:  950 scenarios

u_heighten: [61.4, 250.0]
```

Your Response — Question 2

Look at the PRIM box above. Complete this sentence: "Total cost is unacceptably high when ____ AND ____." Use the parameter names and bounds from the box.

Does the box restrict climate parameters, policy levers, or both? What does that tell you about what drives high cost in this model?

4.7 Input Distributions by Outcome

Now that PRIM has identified a box, let's look directly at the input distributions of passing and failing runs. The histograms below show each sampled variable split by outcome — blue for passing runs, red for failing ones.

```
let
  pass_mask = .!failures
  fail_mask = failures

  fig = Figure(; size=(900, 400))
  positions = [(1, 1), (1, 2), (1, 3), (1, 4)]
  for (j, (name, pos)) in enumerate(zip(param_names, positions))
    row, col = pos
    ax = Axis(fig[row, col]; xlabel=name, ylabel="Density")
    hist!(ax, X[pass_mask, j]; bins=20, normalization=:pdf, color=(:steelblue, 0.6))
    hist!(ax, X[fail_mask, j]; bins=20, normalization=:pdf, color=(:tomato, 0.6))
  end
  Legend(
    fig[2, 1:4],
    [PolyElement(; color=(:steelblue, 0.6)), PolyElement(; color=(:tomato, 0.6))],
    ["Pass (below median cost)", "Fail (above median cost)"];
    orientation=:horizontal,
    tellwidth=false,
  )
)
```

```
fig
end
```




Figure 3: Distribution of each sampled input for passing (blue) and failing (red) runs. Variables that strongly separate the two groups are the primary drivers of high cost.

4.8 Visualization

Your Response — Question 3

Choose any two parameters and create a scatter plot where each point is one scenario, colored by whether it passed or failed. In 2–3 sentences: which parameters best separate the two groups? Does the separation you see match what PRIM identified? Why do some parameters separate the groups while others don't?

We'll share these plots in class — think about which pair of parameters makes the clearest picture.

```
# Your code here
```

5 Key Takeaways

Key takeaways:

1. **Scenario discovery for policy design:** by jointly sampling uncertainties and policy levers, PRIM reveals which combinations drive unacceptable outcomes — not just which conditions, but which decisions under which conditions.
2. **PRIM finds interpretable boxes:** “cost is high when $\eta > X$ AND $u_{\text{heighten}} < Y$ ” is a concrete, actionable statement that connects climate risk to policy choice.
3. **Both levers and uncertainties matter:** PRIM searches the joint space — it can tell you that the policy you control is as important a driver of failure as the uncertainty you can't control.

4. **The structure is visible in the data:** PRIM's box reflects real patterns — the same separation shows up in a scatter plot of raw runs, and is worth examining directly.

6 Submission

1. Write your answers in the response boxes above.
2. **Render to PDF:**
 - In VS Code: Open the command palette (Cmd+Shift+P / Ctrl+Shift+P) → "Quarto: Render Document" → select Typst PDF
 - Or from the terminal: `quarto render index.qmd --to typst`
3. **Submit the PDF** to the Lab 8 assignment on Canvas.
4. **Push your code** to GitHub (for backup):

```
git add -A && git commit -m "Lab 8 complete" && git push
```

Checklist

Before submitting:

- ☐ Packages load without error
- ☐ Joint ensemble runs successfully (1,000 scenario-policy pairs)
- ☐ Cost landscape plotted (Figure 1)
- ☐ Question 1 answered (fixed parameter tweaked and described)
- ☐ PRIM run and peeling trajectory plotted (Figure 2)
- ☐ Question 2 answered (PRIM box interpreted)
- ☐ Pass vs. fail histograms plotted (Figure 3)
- ☐ Question 3 answered (scatter plot created and interpreted)
- ☐ Notebook renders to PDF without errors

Bibliography

- [1] D. van Dantzig, "Economic Decision Problems for Flood Prevention," *Econometrica*, vol. 24, no. 3, pp. 276–287, 1956, doi: 10.2307/1911632.
- [2] C. Eijgenraam *et al.*, "Economically Efficient Standards to Protect the Netherlands Against Flooding," *Interfaces*, vol. 44, no. 1, pp. 7–21, 2014, doi: 10.1287/inte.2013.0721.
- [3] J. H. Friedman and N. I. Fisher, "Bump Hunting in High-Dimensional Data," *Statistics and Computing*, vol. 9, no. 2, pp. 123–143, 1999, doi: 10.1023/A:1008894516817.