

Lab 10: Pareto Fronts

CEVE 421/521

2026-04-10

Warning

Version note: An earlier version of this lab used a house-elevation model instead of the dike model below. If you already started working from that version, that's fine — keep working with the one you have on your computer. Both versions cover the same concepts (Pareto fronts, NSGA-II, static vs. adaptive policies).

1 Overview

Monday's lecture used weighted sums to combine objectives — but weights embed value judgments. What if we showed decision-makers the full trade-off instead?

In this lab you will:

1. Evaluate a **grid of static dike heightenings** and identify which ones are non-dominated on two objectives (investment cost vs. expected damages).
2. Use **NSGA-II** to search for **adaptive** buffer/freeboard policies that decide how much to heighten the dike each year in response to observed water levels.
3. Compare the adaptive Pareto front to the static grid — and to the policy you recommended in Lab 8.

The case study is the **Eijgenraam Ring 15 dike** you met in Lab 8 [1], now reformulated with annual timesteps so adaptive policies can react to observed flood conditions. This follows G. G. Garner and K. Keller [2], who used Direct Policy Search (DPS) to show that adaptive dike heightening strategies can Pareto-dominate static ones.

References:

- C. Eijgenraam *et al.* [1] — the base dike model
- G. G. Garner and K. Keller [2] — Direct Policy Search for sea-level rise adaptation

2 Before Lab

Complete these steps **BEFORE** coming to lab:

1. **Accept the GitHub Classroom assignment** (link on Canvas)
2. **Clone your repository:**

```
git clone https://github.com/CEVE-421-521/lab-10-S26-yourusername.git
cd lab-10-S26-yourusername
```

3. **Open the notebook** in VS Code or Quarto preview
4. **Run the first code cell** — it will install any missing packages (may take 10-15 minutes the first time)

! Important

Submission: Render this notebook to PDF and submit the PDF to Canvas. See Section 5 for details.

3 Setup

Run the cells in this section to load packages and the dike model. You don't need to modify any of this code — just run it.

3.1 Packages

```
# Install SimOptDecisions if needed
if !isfile("Manifest.toml")
    import Pkg
    Pkg.instantiate()
end

try
    using SimOptDecisions
catch
    import Pkg
    try
        Pkg.rm("SimOptDecisions")
    catch
    end
    Pkg.add(; url="https://github.com/dossgollin-lab/SimOptDecisions")
    Pkg.instantiate()
    using SimOptDecisions
end

using CairoMakie
using DataFrames
using Distributions
using Random
using Statistics
using Unitful

import Metaheuristics # loads the optimization extension
```

```
Random.seed!(2026)
```

3.2 The Dike Model

The `eijgenraam_timestep.jl` file defines a timestepping version of the Ring 15 dike model from Lab 8. Each year, a stochastic water level (surge on top of a polynomial SLR trajectory) arrives; the policy decides whether to heighten the dike; failures and costs accumulate with discounting.

💡 Tip

Explore the model. Open `eijgenraam_timestep.jl` in VS Code and skim it. You don't need to memorize every line, but you should be able to identify:

- the **config** (fixed Ring 15 parameters),
- the **scenario** (uncertain inputs: water-level trajectory + economic parameters),
- the two **policies** (static one-shot heightening, adaptive buffer/freeboard rule), and
- the **outcome** metrics (investment cost, expected damages, reliability).

If any of the `SimOptDecisions` macros (`@scenariodef`, `@policydef`, `@outcomedef`) or the callback functions (`get_action`, `run_timestep`, `compute_outcome`) are unfamiliar, ask Copilot / Claude Code / Gemini for a line-by-line walkthrough. The AI assistant can explain the framework in context; you should not be memorizing `SimOptDecisions` syntax.

```
include("eijgenraam_timestep.jl")
```

3.3 Scenarios

We sample 100 scenarios for both parts of the lab.

The only uncertainty we vary is sea-level rise. Each scenario draws a different polynomial SLR trajectory $z_t = a + bt + ct^2$ from broad distributions over the coefficients, and then adds annual storm surges drawn from a fixed GEV calibrated to the Sewells Point, Norfolk VA tide gauge [3]. The starting dike height, the storm surge distribution, the discount rate, and the economic parameters are all **fixed** — students should be able to look at differences across scenarios and attribute them to one thing: is SLR fast or slow?

Sharing the scenario set across policies means differences we observe are **attributable to the policy**, not to sampling noise.

```
config = DikeConfig()
n_scenarios = 100
rng = Random.Xoshiro(2026)
scenarios = [sample_dike_scenario(rng, config) for _ in 1:n_scenarios]

println("Config: $(config.horizon)-year horizon, starting dike H0 = $(config.H0) cm")
println("Scenarios: $(n_scenarios) sampled futures (SLR uncertainty only)")
```

3.4 Surge PDF and SLR Trajectories

Before running any policy, let's look at what we actually sampled: the fixed Norfolk-calibrated GEV surge distribution [3] (converted to feet using `Unitful`) and the polynomial MSL trajectories across 20 representative scenarios.

```
let
    fig = Figure(; size=(900, 400))

    # --- Surge PDF in feet ---
    cm_to_ft(x) = ustrip(u"ft", x * u"cm")
    m_to_ft(x)  = ustrip(u"ft", x * u"m")

    surge_dist_m = GeneralizedExtremeValue(
        config.mu_surge_m, config.sigma_surge_m, config.xi_surge
    )
    xs_m = range(0.0, 4.0; length=400)           # meters
    pdf_m = pdf.(surge_dist_m, xs_m)             # density on meters
    xs_ft = m_to_ft.(xs_m)
    # Convert density: if X in m, Y = X * 3.28084 ft/m, so f_Y(y) = f_X(x)/3.28084
    pdf_ft = pdf_m ./ 3.28084

    ax1 = Axis(
        fig[1, 1];
        xlabel="Annual max storm surge (ft)",
        ylabel="Density",
        title="Norfolk-like GEV ( $\mu$ =$(config.mu_surge_m) m,  $\sigma$ =$(config.sigma_surge_m)
m,  $\xi$ =$(config.xi_surge))",
    )
    lines!(ax1, xs_ft, pdf_ft; color=:steelblue, linewidth=2)
    # Mark some quantiles
    for (q, label) in [(0.5, "p50"), (0.9, "p90"), (0.99, "p99")]
        qv_m = quantile(surge_dist_m, q)
        qv_ft = m_to_ft(qv_m)
        vlins!(ax1, [qv_ft]; color=:gray, linestyle=:dash)
        text!(ax1, qv_ft, maximum(pdf_ft) * 0.9;
            text=label, align=:left, :top, fontsize=10)
    end

    # --- SLR trajectories in feet ---
    years = 1:(config.horizon)
    ax2 = Axis(
        fig[1, 2];
        xlabel="Year",
        ylabel="MSL anomaly (ft)",
        title="Polynomial SLR trajectories (20 scenarios)",
    )
```

```

# NOTE: these 20 MSL trajectories are sampled INDEPENDENTLY (with their own
# seeded RNG) purely for visualization – they are NOT the same draws as the
# 100 scenarios used in the rest of the lab. We can't recover MSL from a
# scenario's water level after the fact (surge is baked in), so we re-draw
# (a, b, c) triples from the same distributions just to show the shape.
display_rng = Random.Xoshiro(7)
for i in 1:20
    a = rand(display_rng, Normal(0.0, 5.0))
    b = rand(display_rng, truncated(Normal(0.3, 0.1); lower=0.05))
    c = rand(display_rng, truncated(Normal(0.005, 0.003); lower=0.0))
    z_cm = [a + b * t + c * t^2 for t in years]
    z_ft = cm_to_ft.(z_cm)
    lines!(ax2, years, z_ft; color=(:steelblue, 0.4), linewidth=1)
end
fig
end

```




Figure 1: Left: the fixed Norfolk-calibrated GEV annual-max storm surge distribution (PDF, in feet). Right: polynomial MSL trajectories $z_t = a + bt + ct^2$ for 20 representative scenarios (feet above baseline).

3.5 Metrics

When the optimizer evaluates a candidate policy, it runs that policy on all 100 scenarios and gets 100 `DikeOutcome` values back — one per scenario. It then needs a single number (actually two, since we have two objectives) per policy to compare against others. That’s what `calculate_metrics` does: it **aggregates outcomes across scenarios** into the two summary statistics we’ll optimize over.

Matching G. G. Garner and K. Keller [2], the two objectives are:

- `total_investment` — average discounted investment cost across scenarios, in M€
- `total_loss` — average discounted expected damages across scenarios, in M€

We also compute `reliability` (fraction of years without overtopping) as a diagnostic alongside the two objectives.

```
function calculate_metrics(outcomes)
    invest = [SimOptDecisions.value(o.investment_cost) for o in outcomes]
    damage = [SimOptDecisions.value(o.expected_damages) for o in outcomes]
    rel     = [SimOptDecisions.value(o.reliability)      for o in outcomes]
    return (
        total_investment = mean(invest),
        total_loss       = mean(damage),
        reliability       = mean(rel),
    )
end
```

3.6 A helper for trade-off plots

All our trade-off plots share the same axes (discounted investment cost vs. discounted expected damages, both in M€), and we'd like them to share the same visual conventions too:

- square aspect (both axes are in the same units),
- axes capped at 10^6 M€ — this is effectively a “no policy costing more than a trillion euros” constraint, used purely to keep the view legible,
- iso-cost reference lines (`investment + damage = constant`), which run top-left to bottom-right and let you read off the total-cost ordering by eye.

```
# All plot data is converted from M€ (the internal unit, matching Eijgenraam
# Table 1) to  $10^9$  € for display. Both axes use a pseudo-log scale so that the
# interesting low-cost region is visible while the overbuilt extremes still
# fit on the plot. Zero values are drawn natively (pseudolog handles zero).
```

```
const ME_PER_BILLION = 1e3      # M€ per  $10^9$  €
```

```
ax_cap_b = 5.0    #  $10^9$  € – visual cap for both axes
```

```
function tradeoff_axis(fig, position; title="", cap=ax_cap_b)
```

```
    ax = Axis(
        fig[position...];
        xlabel="Discounted investment cost ( $10^9$  €)",
        ylabel="Discounted expected damages ( $10^9$  €)",
        title=title,
        aspect=1,
        xscale=Makie.pseudolog10,
        yscale=Makie.pseudolog10,
        limits=(-0.02, cap, -0.02, cap),
        xticks=[0, 0.1, 0.3, 1, 3, 5],
        yticks=[0, 0.1, 0.3, 1, 3, 5],
    )
```

```
    # Iso-cost reference lines: investment + damage = constant.
```

```

# These become curved under pseudolog but still indicate iso-cost levels.
for level in (cap / 5, cap / 2, cap)
    xs = range(0.0, level; length=50)
    ys = level .- xs
    lines!(ax, xs, ys; color=:gray, 0.35), linestyle=:dash, linewidth=1)
end
# Preference star: bottom-left is ideal.
scatter!(ax, [0.0], [0.0];
        marker=:star5, color=:gold, strokecolor=:black, strokewidth=1,
        markersize=22)
return ax
end

"""Convenience: convert a vector of M€ values to 109 €."""
to_billion(v) = v ./ ME_PER_BILLION

```

Static Grid and Trade-offs

Before running any optimizer, let's evaluate a grid of static heightenings from 0 to 4 m (on top of the existing 2.5 m dike) and see the shape of the trade-off space.

3.7 Run the Grid

```

heightenings = collect(0.0:20.0:400.0) # cm, 0–4 m
static_policies = [StaticDikePolicy(; u_heighten=u) for u in heightenings]
static_result = explore(config, scenarios, static_policies; progress=false)

```

Extract per-policy mean metrics and reliability:

```

investment_mat = Array(static_result[:investment_cost])
damage_mat     = Array(static_result[:expected_damages])
reliability_mat = Array(static_result[:reliability])

static_df = DataFrame(;
    u_heighten      = heightenings,
    total_investment = [mean(investment_mat[p, :]) for p in axes(investment_mat, 1)],
    total_loss       = [mean(damage_mat[p, :])      for p in axes(damage_mat, 1)],
    reliability      = [mean(reliability_mat[p, :]) for p in axes(reliability_mat, 1)],
)
static_df

```

i Note

Checkpoint: At `u_heighten = 0`, investment cost should be zero. As `u_heighten` grows, investment climbs (roughly exponentially, from the Eijgenraam cost formula) and damages shrink.

3.8 Non-dominated Set

A static policy is **non-dominated** on the two objectives if no other policy achieves lower investment cost AND lower damages simultaneously.

```
"""
    is_nondominated(i, df, cols) -> Bool

Return `true` if row `i` of `df` is non-dominated on the objective columns `cols`
(all to be minimized). Row `i` is dominated only if some other row `j` is
**weakly better on every objective** and **strictly better on at least one**.
"""

function is_nondominated(i, df, cols)
    row_i = df[i, :]
    for j in axes(df, 1)
        i == j && continue
        row_j = df[j, :]
        all_leq = all(row_j[c] <= row_i[c] for c in cols)
        any_lt = any(row_j[c] < row_i[c] for c in cols)
        all_leq && any_lt && return false
    end
    return true
end

obj_cols = [:total_investment, :total_loss]
mask = [is_nondominated(i, static_df, obj_cols) for i in axes(static_df, 1)]
static_nondom = static_df[mask, :]

println("$(nrow(static_nondom)) of $(nrow(static_df)) static policies are non-
dominated")
```

```
9 of 21 static policies are non-dominated
```

3.9 Plot the Trade-off

```
let
    fig = Figure(; size=(650, 650))
    ax = tradeoff_axis(fig, (1, 1); title="Static Dike Heightening")
    scatter!(
        ax,
        to_billion(static_df.total_investment),
        to_billion(static_df.total_loss);
        color=static_df.u_heighten, colormap=:viridis, markersize=11,
        label="Static (grid)",
    )
    scatter!(
```



```

        ax,
        to_billion(static_nondom.total_investment),
        to_billion(static_nondom.total_loss);
        color=:red, marker=:diamond, markersize=10, label="Non-dominated",
    )
    Colorbar(fig[1, 2]; colormap=:viridis,
        limits=(minimum(static_df.u_heighten), maximum(static_df.u_heighten)),
        label="Heightening (cm)")
    axislegend(ax; position=:rt)
    fig
end

```




Figure 2: Static heightening grid. Each point is a static one-shot heightening; the red diamonds mark the non-dominated subset (the static Pareto front). Dashed curves are iso-cost reference lines (investment + damage = constant; curved because of the pseudo-log axes).

Your Response — Question 1

Look at the static grid plot.

1. *Where does the policy you recommended in Lab 8 fall on this plot? (Re-run the grid including your Lab 8 `u_heighten` value if needed.) Is it on the non-dominated set?*
2. *How many policies are non-dominated? Describe the shape of the trade-off.*
3. *Where is the “knee” of the trade-off — the point where additional investment stops buying much damage reduction?*

Adaptive Policies

The static grid shows the best one-shot heightenings. But a real dike manager can observe rising water levels and decide to heighten later — trading upfront investment for information.

In most settings adaptivity pays off: waiting to see whether SLR is fast or slow before committing to a heightening should beat any one-shot rule. But “should” is not “always” — if the information arrives too late, or the rule you parameterize can’t tell a one-off surge from a sustained trend, a well-chosen static heightening can match or even beat the adaptive front. Keep an eye out for which regime you’re in.

3.10 Background: Garner & Keller (2018)

In G. G. Garner and K. Keller [2], the authors reformulate a Dutch dike-heightening problem to explicitly handle uncertainty in sea-level rise and storm surge, and compare several formulations of the decision problem. Three pieces of their setup are directly relevant to this lab:

- **Two objectives.** Both minimized: discounted investment cost, and discounted expected damages. We use exactly these.
- **Direct Policy Search (DPS) with a buffer/freeboard rule.** Each year the policy looks at the gap between the observed water level and the top of the dike; if the gap is less than a **buffer** height, it raises the dike to restore the buffer plus an extra **freeboard** height. The paper also explores a richer 10-variable version of the DPS rule that lets the buffer and freeboard depend on observed SLR trend and variability — but it reports that the simple 2-variable constant rule already captures most of the gain, and the 10-variable version adds only a modest improvement on top. Our lab uses the 2-variable version.
- **Borg MOEA.** The paper uses a parallel Borg Multi-Objective Evolutionary Algorithm, run for 50 independent seeds $\times$ 200,000 function evaluations each (= 10 million evaluations). Borg combines auto-adaptive operator selection, restart logic, and an archive that explicitly controls objective-space spacing on top of a plain genetic algorithm. We don’t have Borg available in Julia, so we use **NSGA-II** from `Metaheuristics.jl` — a well-established multi-objective genetic algorithm. See the NSGA-II Wikipedia entry for a short description of how it works.

i Optimizers don't give you a free lunch

NSGA-II can and will produce a front that looks *uneven* — dense in some regions and sparse in others — even when the population is “converged”. The reason is that the genetic operators (crossover + mutation) work in **decision space**, while the diversity operator (crowding distance) works in **objective space**; if the decision-space-to-objective-space mapping is strongly non-linear (as it is here, because the Eijgenraam cost is roughly exponential in heightening), uniform exploration in decision space produces very non-uniform coverage of the Pareto front.

The standard tools for handling this are, in order of increasing effort:

1. **Tighten decision-variable bounds** so that the search doesn't waste effort on regions that are dominated by construction.
2. **Bigger populations** so crowding distance has more points to spread around.
3. **Reparameterize** the decision variables (e.g. work in log-cost space, or redefine the policy in terms of target dike height rather than buffer cm).
4. **A different algorithm**: `Metaheuristics.jl` also ships `SMS-EMOA`, `SPEA2`, and `MOEA/D-DE`, each with a different diversity mechanism. Borg's ϵ -dominated archive gives explicit control over objective-space spacing.
5. **Multiple random seeds** and merging the archives — the approach G. G. Garner and K. Keller [2] took with 50 seeds.

We've already done some of (1) for this lab — the adaptive bounds are tuned so that most of the sampled decision space lands in useful objective-space territory. You can see residual non-uniformity in the resulting Pareto front.

What we simplify relative to GK18:

Element	Garner & Keller (2018)	This lab
Horizon	300 years	100 years
Scenarios	100,000	100
DPS form	10 variables (state-dependent)	2 variables (constant buffer/free-board)
Optimizer	Borg MOEA, 50 seeds $\times$ 200k evals	NSGA-II, 1 seed, $\sim$ 100 iterations

3.11 Brute-Force Grid Over the Adaptive Policy Space

Before turning to NSGA-II, let's do the same thing we did for static policies: evaluate a grid of (`buffer`, `freeboard`) combinations and see the trade-off directly. With only two decision variables, a modest grid gives us a clear picture of what the *true* Pareto front looks like, and gives us something to compare the optimizer against.

```
buffers      = collect(0.0:25.0:400.0)  # cm, 0–4 m
freeboards   = collect(0.0:25.0:300.0)  # cm, 0–3 m
```

```

adaptive_grid = [
    (buffer=buf, freeboard=fb, policy=AdaptiveDikePolicy(; buffer=buf, freeboard=fb))
    for buf in buffers, fb in freeboards
]
adaptive_grid_vec = vec(adaptive_grid)

adaptive_result = explore(
    config, scenarios,
    [g.policy for g in adaptive_grid_vec];
    progress=false,
)

adaptive_inv_mat = Array(adaptive_result[:investment_cost])
adaptive_dmg_mat = Array(adaptive_result[:expected_damages])
adaptive_rel_mat = Array(adaptive_result[:reliability])

adaptive_grid_df = DataFrame(;
    buffer      = [g.buffer      for g in adaptive_grid_vec],
    freeboard   = [g.freeboard   for g in adaptive_grid_vec],
    total_investment = [mean(adaptive_inv_mat[p, :]) for p in axes(adaptive_inv_mat, 1)],
    total_loss    = [mean(adaptive_dmg_mat[p, :]) for p in axes(adaptive_dmg_mat, 1)],
    reliability    = [mean(adaptive_rel_mat[p, :]) for p in axes(adaptive_rel_mat, 1)],
)

grid_mask      = [is_nondominated(i, adaptive_grid_df, obj_cols) for i in
axes(adaptive_grid_df, 1)]
adaptive_grid_nondom = adaptive_grid_df[grid_mask, :]

```

```

println("$ (nrow(adaptive_grid_nondom)) of $ (nrow(adaptive_grid_df)) grid policies are
non-dominated")

```

16 of 221 grid policies are non-dominated

```

let
    fig = Figure(; size=(650, 650))
    ax = tradeoff_axis(fig, (1, 1); title="Adaptive Policy Grid – Buffer/Freeboard")
    scatter!(
        ax,
        to_billion(static_nondom.total_investment),
        to_billion(static_nondom.total_loss);
        color=:darkorange, marker=:utriangle, markersize=11,
        label="Static non-dominated",
    )

```

```

scatter!(
    ax,
    to_billion(adaptive_grid_df.total_investment),
    to_billion(adaptive_grid_df.total_loss);
    color=(steelblue, 0.35), markersize=6,
    label="Adaptive (grid)",
)
scatter!(
    ax,
    to_billion(adaptive_grid_nondom.total_investment),
    to_billion(adaptive_grid_nondom.total_loss);
    color=:red, marker=:diamond, markersize=10,
    label="Grid non-dominated",
)
axislegend(ax; position=:rt)
fig
end

```




Figure 3: Brute-force grid over the (buffer, freeboard) policy space. Each point is one adaptive policy; red diamonds mark the non-dominated set. The static non-dominated points are shown in orange for reference.

The grid gives us a **visual ground truth** for what the adaptive Pareto front looks like. In the next section we hand the same problem to NSGA-II and check whether it recovers the same shape.

3.12 Configure the Optimizer

SimOptDecisions can dispatch to several optimization backends. Here we use the `MetaheuristicsBackend`, which wraps `Metaheuristics.jl` and lets us pick an algorithm (`:NSGA2`), a population size, and a number of iterations. We also set `store_history = true` so that, at the end, we can look at how the Pareto front evolved across generations — not just the final result.

```

backend = MetaheuristicsBackend(;
    algorithm      = :NSGA2,      # multi-objective GA
    max_iterations = 100,         # number of generations (NSGA-II may early-stop
    if converged)
    population_size = 60,         # policies carried forward each generation
    parallel        = true,       # evaluate scenarios in parallel where possible
    store_history    = true,       # keep the per-generation archive for plotting
)

```

Each generation, NSGA-II evaluates 60 candidate policies, each on all 100 scenarios (so 6,000 simulations per generation). NSGA-II typically early-stops around 50 generations once the population stops improving. That should run in well under a minute.

Default objectives matching GK18: **minimize investment cost and expected damages**. You can swap these for other pairs (e.g. include `reliability` as a third objective), but each re-run takes a minute or two.

```
# Your code here
```

3.13 Run Optimization

We time the optimization so you can see how long it took and plan accordingly for the final run.

```
# Your code here
```

i Note

This is a first run with 100 scenarios and 50 iterations to keep things fast. **Before you write up your conclusions**, re-run this block (and the plots below) with more scenarios (e.g. 500) and more iterations (e.g. 100) so the Pareto front is stable. Expect ~5 minutes for the final run.

3.14 The Final Pareto Front

Plot the adaptive Pareto front on top of the static grid for comparison.

```
# Your code here
```

Figure 4

3.15 Watching the Front Converge

Because we set `store_history = true`, we can extract the Pareto archive at every generation and watch it evolve. Early generations are noisy; later ones should concentrate near the final front.

```
# Your code here
```

Figure 5

Your Response — Question 2

Pick one solution from the adaptive Pareto front.

1. *What buffer and freeboard values does it use?*
2. *Describe in plain language what this policy does — when does it trigger heightening, and how much does it raise the dike?*
3. *Who might prefer this solution over the alternatives on the Pareto front — a risk-averse stakeholder, a fiscally-conservative one, or someone in the middle? Why?*

Your Response — Question 3

Compare the adaptive Pareto front to the static grid and to your Lab 8 recommendation.

1. *Does the adaptive front Pareto-dominate all of the static policies? Some of them? None?*
2. *What is the source of the adaptive policy's advantage (when it has one)?*
3. *GK18 also studied a 10-variable DPS that lets the buffer and freeboard respond to the observed SLR trend and variability. In their results, the 2-variable version captured most of the gain. Does that match your intuition from this lab — and under what conditions would you expect the 10-variable version to matter more?*

4 Wrap-Up

Key takeaways:

1. **Choosing objectives is a value judgment.** The Pareto front changes depending on which metrics you optimize. The analyst's job is to present trade-offs, not prescribe a single answer.
2. **The Pareto front separates analysis from decision.** It shows what is technically achievable; stakeholders with different risk tolerances, budgets, and values can each identify their preferred solution.
3. **Adaptive policies can dominate static ones.** By conditioning actions on observed conditions, adaptive strategies exploit information revealed over time.
4. **Algorithm choices matter.** NSGA-II is a solid workhorse for 2–3 objective problems, but real MORDM work often uses Borg MOEA with ϵ -dominance and many random seeds. Our classroom setup is a simplified version of what G. G. Garner and K. Keller [2] used.

5 Submission

1. Write your answers in the response boxes above.
2. **Before writing up**, re-run the optimizer with more scenarios and more iterations (see the callout above) so your Pareto front is stable.
3. **Render to PDF:**
 - In VS Code: command palette (`Cmd+Shift+P` / `Ctrl+Shift+P`) → “Quarto: Render Document” → select Typst PDF
 - Or from the terminal: `quarto render index.qmd --to typst`

4. **Submit the PDF** to the Lab 10 assignment on Canvas.
5. **Push your code** to GitHub (for backup):

```
git add -A && git commit -m "Lab 10 complete" && git push
```

Checklist

Before submitting:

- ☐ Packages load without error
- ☐ Setup cells run (model, scenarios, metrics)
- ☐ Static grid evaluated with `explore()`
- ☐ Static non-dominated set identified and plotted (Figure 2)
- ☐ Question 1 answered (Lab 8 comparison, shape of trade-off, reliability filter)
- ☐ NSGA-II optimization run with `store_history=true`
- ☐ Final Pareto front plotted over static grid (Figure 4)
- ☐ Pareto-front evolution plotted across generations (Figure 5)
- ☐ Question 2 answered (interpret one adaptive solution)
- ☐ Question 3 answered (compare adaptive vs. static, reflect on 2- vs. 10-variable DPS)
- ☐ Final run completed with more scenarios / iterations
- ☐ Notebook renders to PDF without errors

Bibliography

- [1] C. Eijgenraam *et al.*, “Economically Efficient Standards to Protect the Netherlands Against Flooding,” *Interfaces*, vol. 44, no. 1, pp. 7–21, 2014, doi: 10.1287/inte.2013.0721.
- [2] G. G. Garner and K. Keller, “Using Direct Policy Search to Identify Robust Strategies in Adapting to Uncertain Sea-Level Rise and Storm Surge,” *Environmental Modelling & Software*, vol. 107, pp. 96–104, 2018, doi: 10.1016/j.envsoft.2018.05.006.
- [3] J. Doss-Gollin and K. Keller, “A Subjective Bayesian Framework for Synthesizing Deep Uncertainties in Climate Risk Management,” *Earth’s Future*, vol. 11, no. 1, Jan. 2023, doi: 10.1029/2022EF003044.